

Turing quickstart Beta-Bernoulli

```
1 #     margin-left: 1%;  
2 #     margin-right: 5%;  
3 html"""<style>  
4 main {  
5     margin: 0 auto;  
6     max-width: 90%;  
7     padding-left: max(50px, 1%);  
8     padding-right: max(253px, 10%);  
9     # 253px to accomodate TableOfContents(aside=true)  
10 }  
11 """"
```

```
1 begin  
2     using Markdown  
3     using InteractiveUtils  
4     using Pkg, DrWatson, PlutoUI  
5 end
```

```
► TaskLocalRNG()  
1 begin  
2     using Distributions  
3     using FillArrays  
4     using StatsPlots  
5  
6     using LinearAlgebra  
7     using Random  
8     using Turing  
9  
10    # Import all libraries.  
11    using DataFrames  
12    #DirichletProcess, ChineseRestaurantProcess,  
13    StickBreakingProcess  
14    using Turing.RandomMeasures  
15    # DocStringExtensions provides $(SIGNATURES)  
16    using DocStringExtensions  
17    # Set a random seed.  
18    Random.seed!(3)  
end
```

```
1 using DynamicPPL, Printf
```

☰ Table of Contents

Turing quickstart Beta-Bernoulli

1. Turing.jl

- 1.1 Conditioning and Deconditioning
 - 1.1.1 `model1_1`: `gdemo` conditioned on `dataset1_1` and $\mu=0$
 - 1.1.2 Correct syntax to define a conditional distribution with a Named...
 - 1.1.3 Sample from a fake conditional distribution (a full distribution in...
 - 1.1.4 Sample μ from `gdemo()` conditioned on an x .
 - 1.1.5 Sample x from `gdemo()` conditioned on an $\mu=15$.
 - 1.1.6 `gdemo2()`: $x(\text{observations})$ as input argument

2. The Problem: Coin flipping

- 2.1 Question
- 2.2 Input: flip 100 times

3 Frequentist/Likelihood approach

- 3.1 The MLE (maximum likelihood estimator) is $\hat{p} = \sum_{i=1}^{i=N} x_i/N$ and why?

4 Bayesian

- 4.1 Declare our Bayesian model and Monte Carlo with a HMC sampler
 - 4.1.1 The posterior of p is max at $p=0.4311$, almost same as MLE. No...
- 4.2 Input is 10X more data.
 - 4.2.1 sample using the same model but with 10X iterations
 - 4.2.2 same model, but using the NUTS sampler (default parameters)
- 4.3 Modify the prior to alter the posterior
 - 4.3.1 Conclusion

5 Test two different Negative Binomial models

- 5.1 NB with a beta prior for p and Uniform prior for r
- 5.2 NB with a beta prior for p and Gamma prior for r

6 Estimate p of Binomial(N, p) from Binomial data with different Ns

- 6.1 Estimate if true p is 0.0005
- 6.2 Estimate if true p is 0.005
- 6.3 Estimate if true p is 0.00001

7 Estimate p of NegativeBinomial(r, p) from Binomial data with differe...

- 7.1 \hat{p} of Beta+Binomial is similar to \hat{p} of Beta+NegativeBinomial.

8 Multinomial

10 Visualize beta distributions

```
1 begin
2   PlutoUI.TableOfContents()
3 end
```

```
1 versioninfo()

[+] Julia Version 1.10.2
Commit bd47eca2c8a (2024-03-01 10:14 UTC)
Build Info:
  Official https://julialang.org/ release
Platform Info:
  OS: Linux (x86_64-linux-gnu)
  CPU: 32 × Intel(R) Xeon(R) CPU E5-2630 v3 @ 2.40GHz
  WORD_SIZE: 64
  LIBM: libopenlibm
  LLVM: libLLVM-15.0.7 (ORCJIT, haswell)
Threads: 16 default, 0 interactive, 8 GC (on 32 virtual cores)
Environment:
  JULIA_PKG_SERVER = https://mirrors.tuna.tsinghua.edu.cn/julia
  JULIA_REVISE_WORKER_ONLY = 1
```

```
"foo bar 34.567"
1 begin
2   x = "foo"
3   println("$x $(x)", "haha")
4   #using Printf
5   @sprintf "%s %s %.3f" "foo" "bar" 34.567
6 end
```

```
foo foo haha
```

1. Turing.jl

- <https://turinglang.org/>
- <https://turinglang.org/DynamicPPL.jl/stable/>
- <https://github.com/TuringLang/Turing.jl>

```
gdemo (generic function with 2 methods)
1 @model function gdemo(n)
2   μ ~ Normal(0, 1)
3   x ~ MvNormal(Fill(μ, n), I)
4 end

▶ [0.848878, -0.319361, -1.39821, -0.0519893, -1.14651, -0.630617,
  ◀ ▶
1 begin
2   Random.seed!(1776)
3   dataset1_1 = randn(100)
4 end

▶ [0.71737]
1 Random.seed!(16); dataset1_2 = randn(1)
```

1.1 Conditioning and Deconditioning

Bayesian models can be transformed with two main operations, conditioning and deconditioning (also known as marginalization).

Conditioning takes a variable and fixes its value as known. We do this by passing a model and a collection of conditioned variables to `|` or its alias `condition`.

1.1.1 model1_1: gdemo conditioned on dataset1_1 and $\mu=0$

- This conditional distribution contains no random variables to draw.
Drawing from it returns empty.
- This model contains both data and parameter values.

```
1 md"## 1.1.1 `model1_1`: gdemo conditioned on `dataset1_1`"
2 and μ=0
3
4   - This conditional distribution contains no random
  variables to draw. Drawing from it returns empty.
  - This model contains both data and parameter values."
```

```
m1_1 =
▶ Model(gdemo (generic function with 2 methods), (n = 100), (), Cc
  ◀ ▶
1 m1_1 = gdemo(length(dataset1_1)) | (x=dataset1_1, μ=0)
```

This operation can be reversed by applying decondition

```
► Model(gdemo (generic function with 2 methods), (n = 100), (), De
◀ ────────────────────────────────────────────────────────────────── ▶
1 DynamicPPL.decondition(m1_1)
```

- Drawing from it returns empty

```
► ()
1 rand(m1_1)
```

We often want to calculate the (unnormalized) probability density for an event. This probability might be a prior, a likelihood, or a posterior (joint) density. DynamicPPL provides convenient functions for this. For example, we can calculate the joint probability of a set of samples (here drawn from the prior) with logjoint:

```
-151.5171285609121
1 begin
2 Random.seed!(124)
3 sample1_1 = rand(m1_1)
4 logjoint(m1_1, sample1_1)
5 end
```

```
► ()
1 sample1_1
```

For models with many variables `rand(model)` can be prohibitively slow since it returns a `NamedTuple` of samples from the prior distribution of the unconditioned variables. We recommend working with samples of type `DataStructures.OrderedDict` in this case:

```
-151.5171285609121
1 begin
2     using DataStructures
3
4     Random.seed!(124)
5     @show sample_dict = rand(OrderedDict, m1_1)
6     logjoint(m1_1, sample_dict)
7 end
```

```
sample_dict = rand(OrderedDict, m1_1) = OrderedCollections.OrderedDict{Any, Any}()
```

```
► OrderedDict()
1 sample_dict
```

- If `sample1_1/sample_dict` is empty, what are `logjoint` and etc. calculating?
- It is computing the joint of the data & parameters already embedded in the model.

```
1 md"- If sample1_1/sample_dict is empty, what are logjoint
2 and etc. calculating?
- It is computing the joint of the data & parameters
already embedded in the model."
```

```
true
1 begin
2   @show logjoint(m1_1, sample1_1)
3   @show loglikelihood(m1_1, sample1_1)
4   @show logprior(m1_1, sample1_1)
5   logjoint(m1_1, sample1_1) ≈ loglikelihood(m1_1,
6   sample1_1) + logprior(m1_1, sample1_1)
end
```

```
logjoint(m1_1, sample1_1) = -151.5171285609121
loglikelihood(m1_1, sample1_1) = -151.5171285609121
logprior(m1_1, sample1_1) = 0.0
```

-151.5171285609121

```
1 @time logjoint(m1_1, sample_dict)
```

```
0.000006 seconds (4 allocations: 1.109 KiB)
```

m1_1_1 =

```
► Model(gdemo (generic function with 2 methods), (n = 1), (), Conc
◀ ────────────────── ▶
1 m1_1_1 = gdemo(length(dataset1_2)) | (x=dataset1_2, μ=15)
```

0.0

```
1 begin
2   @show s1_1_1 = rand(m1_1_1)
3   @show logjoint(m1_1_1, s1_1_1)
4   @show loglikelihood(m1_1_1, s1_1_1)
5   @show logprior(m1_1_1, s1_1_1)
6 end
```

```
s1_1_1 = rand(m1_1_1) = NamedTuple()
logjoint(m1_1_1, s1_1_1) = -216.33463747881498
loglikelihood(m1_1_1, s1_1_1) = -216.33463747881498
logprior(m1_1_1, s1_1_1) = 0.0
```

Sometimes it is helpful to define convenience functions for conditioning on some variable(s). For instance, in this example we might want to define a version of gdemo that conditions on one observation of x:

```
gdemo_x (generic function with 1 method)
1 # this function is essentially same as gdemo2.
2 gdemo_x(x::AbstractVector{<:Real}) = gdemo(length(x)) | (;
  x)
```

1.1.2 Correct syntax to define a conditional distribution with a NamedTuple

```
► (μ = -0.668001)
1 Random.seed!(124); rand(gdemo(length(dataset1_1)) |
(x=dataset1_1,))
```

```
► (μ = -0.668001)
1 Random.seed!(124); rand(gdemo(length(dataset1_1)) |
(;x=dataset1_1,))
```

```
► (μ = -0.668001)
1 Random.seed!(124); rand(gdemo(length(dataset1_1)) |
(;x=dataset1_1))
```

Following syntax is wrong in defining a conditional distribution.

- `(x=dataset1_1)` is not a NamedTuple. It defines a variable named `x`.

```
Random.seed!(124) ; rand(gdemo(length(dataset1_1)) | (x=dataset1_1) )
```

1.1.3 Sample from a fake conditional distribution (a full distribution in fact)

The following is NOT a conditional distribution as `dataset1_1` is not a symbol in the model. It effectively is the full distribution. Check the sampling result in the next block.

```
m1_3 =
▶ Model(gdemo (generic function with 2 methods), (n = 100), (), Conc
◀ └── 1 m1_3 = gdemo(length(dataset1_1)) | (; dataset1_1)

▶ (μ = -0.389096, x = [-0.805381, -2.01082, -0.583562, -0.788156,
◀ └── 1 rand(m1_3)
```

1.1.4 Sample μ from `gdemo()` conditioned on an x .

- This conditional distribution of μ is NOT μ posterior, but the prior distribution.
- Evidenced from the μ histogram.

```
▶ [0.71737]
1 dataset1_2

m1_4 =
▶ Model(gdemo (generic function with 2 methods), (n = 1), (), Conc
◀ └── 1 m1_4 = gdemo(length(dataset1_2)) | (; x=dataset1_2 )

▶ OrderedDict(μ => 0.0419453)
1 @time rand(OrderedDict, m1_4)

[?] 0.021545 seconds (35.39 k allocations: 2.442 MiB, 99.55% compilation time)
```

```
▶ (μ = -2.28973)
1 @time rand(m1_4)

[?] 0.000061 seconds (36 allocations: 2.547 KiB)
```

```
-0.22876723439327667
1 rand(OrderedDict, m1_4).vals[1]

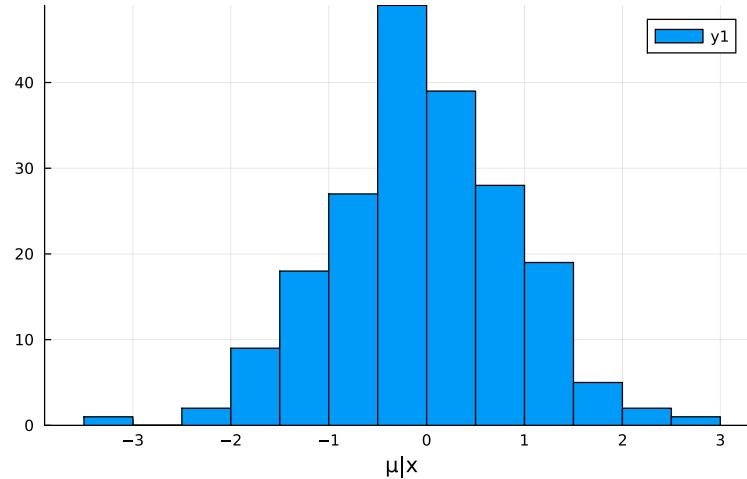
n_samples = 200
1 n_samples = 200
```

```
► [0.540481, -1.0629, -0.729338, 0.259889, 0.261761, 0.555702, 0.0
```

```
1 @time begin
2     # draw  $\mu$  repeatedly and draw a histogram
3     μ_posterior_vec = Vector{Float64}(undef, n_samples)
4     for i ∈ 1:n_samples
5         sample_dict = rand(OrderedDict, m1_4)
6         μ_posterior_vec[i] = sample_dict.vals[1]
7     end
8     μ_posterior_vec
9 end
```

```
0.004735 seconds (6.22 k allocations: 455.969 KiB, 81.21% compilation time)
```

x=[0.7173699612147298]



```
1 histogram(μ_posterior_vec, xlabel="μ|x",
2             title="x=$(dataset1_2)")
```

1.1.5 Sample x from gdemo() conditioned on an $\mu=15$.

- This truly is a conditional distribution, evidenced by the histogram of sampled x

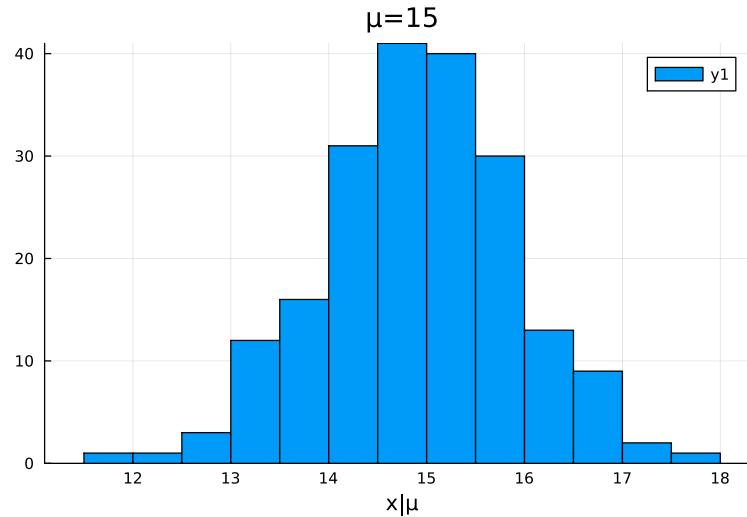
```
m1_5 =
► Model(gdemo (generic function with 2 methods), (n = 1), (), Conc
1 m1_5 = gdemo(1) | (;μ=15)
```

```
16.318092777878245
1 rand(OrderedDict, m1_5).vals[1][1]
```

```
[14.8499, 13.4372, 14.3674, 13.5304, 15.7144, 14.9822, 15.2636, :
```

```
1 @time begin
2     # draw repeatedly and draw a histogram
3     x_posterior_vec = Vector{Float64}(undef, n_samples)
4     for i ∈ 1:n_samples
5         sample_dict = rand(OrderedDict, m1_5)
6         x_posterior_vec[i] = sample_dict.vals[1][1]
7     end
8     x_posterior_vec
9 end
```

```
0.0000913 seconds (6.20 k allocations: 473.641 KiB)
```



```
1 histogram(x_posterior_vec, xlabel="x|μ", title="μ=15")
```

- likelihood and prior are **reversed** due to x not being the argument of gdemo()

```
true
```

```
1 begin
2     @show logjoint(m1_5, (;x=dataset1_2))
3     @show loglikelihood(m1_5, (;x=dataset1_2))
4     @show logprior(m1_5, (x=dataset1_2,))
5     logjoint(m1_5, (;x=dataset1_2)) ≈ loglikelihood(m1_5,
6     (;x=dataset1_2)) + logprior(m1_5, (x=dataset1_2,))
6 end
```

```
logjoint(m1_5, (; x = dataset1_2)) = -216.334637478
81498
loglikelihood(m1_5, (; x = dataset1_2)) = -113.4189385
3320467
logprior(m1_5, (x = dataset1_2,)) = -102.9156989456103
2
```

- likelihood of x given $\mu=15$

```
-102.92098303320466
```

```
1 logpdf(Normal(15,1), 0.717)
```

- prior of $\mu=15$ given $\mu \sim N(0,1)$

```
-113.41893853320467
1 logpdf(Normal(0,1), 15)
```

```
-216.33463747881498
```

```
1 let
2     @show logjoint(gdemo(1), (;x=dataset1_2, μ=15))
3     @show loglikelihood(gdemo(1), (;x=dataset1_2, μ=15))
4     @show logprior(gdemo(1), (x=dataset1_2, μ=15))
5 end
```

```
logjoint(gdemo(1), (; x = dataset1_2, μ = 15)) = -2 ②
16.33463747881498
loglikelihood(gdemo(1), (; x = dataset1_2, μ = 15)) =
0.0
logprior(gdemo(1), (x = dataset1_2, μ = 15)) = -216.33
463747881498
```

- Likelihood=0!
- Prior = Joint = Likelihood of x + prior of μ .

1.1.6 gdemo2(): x(observations) as input argument

```
gdemo2 (generic function with 2 methods)
1 @model function gdemo2(x)
2     μ ~ Normal(0, 1)
3     x ~ MvNormal(Fill(μ, length(x)), I)
4 end
```

```
-113.41893853320467
```

```
1 let
2     @show logjoint(gdemo2([0.717]), (;x=dataset1_2, μ=15))
3     @show loglikelihood(gdemo2([0.717]), (;x=dataset1_2,
4     μ=15))
5     @show logprior(gdemo2([0.717]), (x=dataset1_2, μ=15))
end
```

```
logjoint(gdemo2([0.717]), (; x = dataset1_2, μ = 1 ②
5)) = -216.33992156640932
loglikelihood(gdemo2([0.717]), (; x = dataset1_2, μ =
15)) = -102.92098303320466
logprior(gdemo2([0.717]), (x = dataset1_2, μ = 15)) =
-113.41893853320467
```

- Likelihood and prior are now correct.

```
-113.41893853320467
```

```
1 let
2     @show logjoint(gdemo2([0.717]), (;μ=15))
3     @show loglikelihood(gdemo2([0.717]), (; μ=15))
4     @show logprior(gdemo2([0.717]), ( μ=15,))
5 end
```

```
logjoint(gdemo2([0.717]), (; μ = 15)) = -216.339921 ②
56640932
loglikelihood(gdemo2([0.717]), (; μ = 15)) = -102.9209
8303320466
logprior(gdemo2([0.717]), (μ = 15,)) = -113.4189385332
0467
```

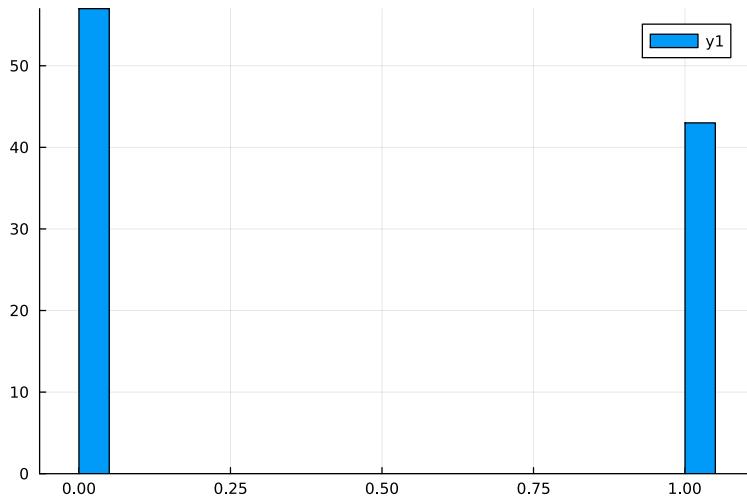
2. The Problem: Coin flipping

Flip a coin. A random variable $x = 1$ if the coin is head up. $x=0$ if the coin is tail up.

2.1 Question

- Input: a sequence of observed x .
- Output: what is the $p(\text{head up})$ of the coin?

2.2 Input: flip 100 times



```
1 begin
2     # Set the true probability of heads in a coin.
3     p_true = 0.5
4
5     # Iterate from having seen 0 observations to 100
6     # observations.
7     Ns = 0:100
8
9     # Fix the seed and draw data from a Bernoulli
10    # distribution, i.e. draw heads or tails.
11    Random.seed!(12)
12    data = rand(Bernoulli(p_true), last(Ns))
13
14    println("The number of ups is
15        $(sum(data)) / $(length(data)).");
16    histogram(data, bins=20)
end
```



The number of ups is 43/100.



3 Frequentist/Likelihood approach

p is an unknown parameter to be solved.

$x_i \sim \text{Bernoulli}(p)$

The MLE (maximum likelihood estimator) is $\hat{p} = \sum_{i=1}^{i=N} x_i/N$ and why?

- What is the difference between likelihood and probability? What is MLE?
- Why most estimators we know/use are MLE estimators?
- Analytic/closed-form (解析) solution vs Monte Carlo solution vs Brute force solution?
 - https://en.wikipedia.org/wiki/Closed-form_expression

```
1 md"# 3 Frequentist/Likelihood approach
2
3 p is an unknown parameter to be solved.
4
5 xi ~ Bernoulli(p)
6
7 The MLE (maximum likelihood estimator) is  $\hat{p} = \sum_{i=1}^{i=N} x_i/N$ 
8 and why?
9
10 - What is the difference between likelihood and
11 probability? What is MLE?
12 - Why most estimators we know/use are MLE estimators?
- Analytic/closed-form (解析) solution vs Monte Carlo
solution vs Brute force solution?
- https://en.wikipedia.org/wiki/Closed-form\_expression"
```

For any model (Likelihood or Bayesian):

- Is there an analytic solution? No for most real-world problems.
- Is there a monte carlo solution? Yes for most real-world problems.
- Is there a brute-force solution? Yes. But may take hundreds of years.

3.1 The MLE (maximum likelihood estimator) is $\hat{p} = \sum_{i=1}^{i=N} x_i / N$ and why?

- <https://stats.stackexchange.com/questions/275380/maximum-likelihood-estimation-for-bernoulli-distribution>

image.png

```
1 md"## 3.1 The MLE (maximum likelihood estimator) is ̂p =  
2 ∑i=1i=N xi/N and why?  
3  
4  
- https://stats.stackexchange.com/questions/275380/maximum-likelihood-estimation-for-bernoulli-distribution  
5  
6 ! [image.png](attachment:f0a44559-d8a5-4842-8ffa-fb2a6d31481f.png)"
```

4 Bayesian

p is a random variable with a prior distribution.

$$p \sim Beta(\alpha, \beta)$$

$$x_i \sim Bernoulli(p)$$

Note: x is Bernoulli, 0 or 1, not Binomial.

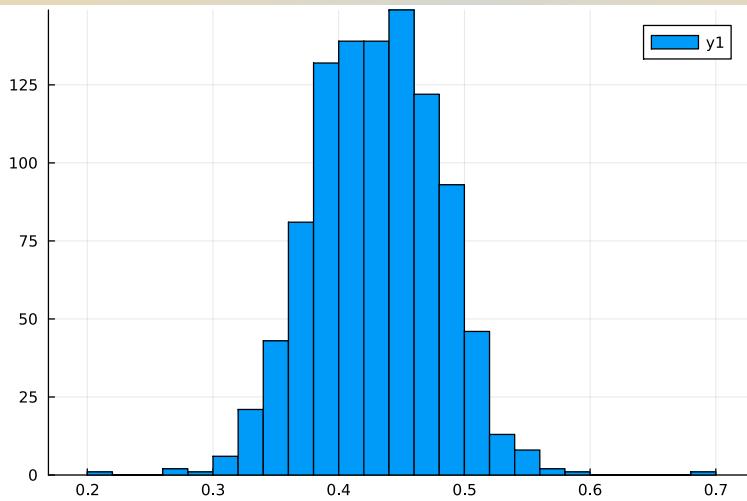
Question: derive the posterior distribution of p , AKA (Also Known As), $\text{Prob}(p|x_1, \dots, x_N)$.

```
1 md"># 4 Bayesian
2
3 p is a random variable with a prior distribution.
4
5 $p \sim Beta(\alpha, \beta)$
6
7 $x_i \sim Bernoulli(p)$
8
9 Note: x is Bernoulli, 0 or 1, not Binomial.
10
11 Question: derive the posterior distribution of $p$, AKA
(Also Known As), $\text{Prob}(p|x_1, \dots, x_N)$."
```

- Is there an analytic solution to this? Yes, Beta-Binomial.
 - https://en.wikipedia.org/wiki/Beta-binomial_distribution
- Is there a monte carlo solution? Yes, we will talk about this today.

4.1 Declare our Bayesian model and Monte Carlo with a HMC sampler

```
coinflip (generic function with 2 methods)
1 # Declare our Bayesian model, p ~ beta(1,1), x ~ Bernoilli(p)
2 @model function coinflip(x)
3     # Our prior belief about the probability of heads in a
4     coin.
5     p ~ Beta(1, 1)
6
7     # The number of observations.
8     N = length(x)
9     for n in 1:N
10        # Heads or tails of a coin are drawn from a
11        Bernoulli distribution.
12        x[n] ~ Bernoulli(p)
13    end
14
```



```

1 begin
2   # Settings of the Hamiltonian Monte Carlo (HMC) sampler.
3   iterations = 1000
4   ε = 0.05
5   τ = 10
6
7   # Start sampling.
8   @time chain = sample(coinflip(data), HMC(ε, τ),
9   iterations)
10  display(chain)
11  # Plot a summary of the sampling process for the
12  # parameter p, i.e. the probability of heads in a coin.
13  histogram(chain[:p])
end

```

100%

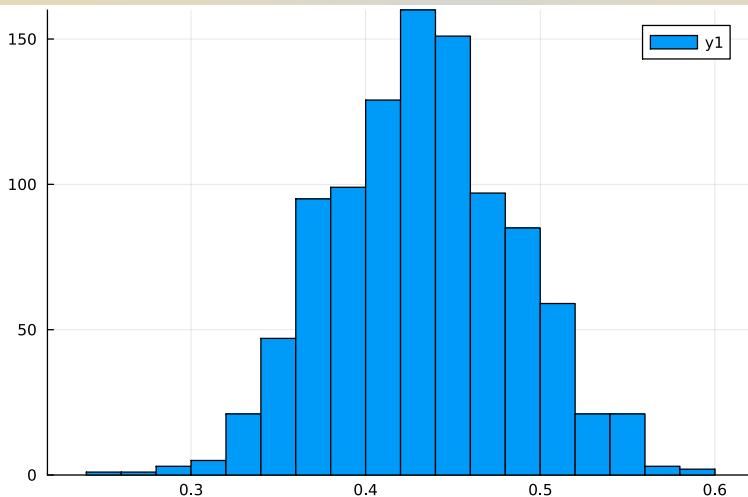
```

13.902008 seconds (9.63 M allocations: 653.702 MiB, 2.78% gc time, 64.55% compilation time)
Chains MCMC chain (1000×11×1 Array{Float64, 3}):
Iterations      = 1:1:1000
Number of chains = 1
Samples per chain = 1000
Wall duration    = 10.75 seconds
Compute duration = 10.75 seconds
parameters       = p
internals        = lp, n_steps, is_accept, acceptance_rate, log_density, hamiltonian_energy, hamiltonian_energy_error, numerical_error, step_size, nom_step_size

Summary Statistics
  parameters      mean       std      mcse    ess_bulk
ess_tail      rh ...
Symbol  Float64  Float64  Float64  Float64
Float64  Float64 ...
421.4425    0.99 ...
p          0.4303   0.0499   0.0009  3000.0000
2 columns omitted

Quantiles
  parameters    2.5%     25.0%    50.0%    75.0%
97.5%
Symbol  Float64  Float64  Float64  Float64
Float64 ...
p        0.3377   0.3950   0.4304   0.4668
0.5197

```



```

1 begin
2   # Using NUTS as sampler
3   @time chain_NUTS = sample(coinflip(data), NUTS(),
4   iterations)
5   display(chain_NUTS)
6   # Plot a summary of the sampling process for the
7   # parameter p, i.e. the probability of heads in a coin.
7   histogram(chain_NUTS[:p])
end

```

100%

ⓘ Found initial step size
ε: 0.8

```

6.243953 seconds (4.15 M allocations: 283.049 MiB, 1.66% gc time, 56.34% compilation time)
Chains MCMC chain (1000×13×1 Array{Float64, 3}):

Iterations      = 501:1:1500
Number of chains = 1
Samples per chain = 1000
Wall duration    = 5.16 seconds
Compute duration = 5.16 seconds
parameters       = p
internals        = lp, n_steps, is_accept, acceptance_rate, log_density, hamiltonian_energy, hamiltonian_energy_error, max_hamiltonian_energy_error, tree_depth, numerical_error, step_size, nom_step_size

Summary Statistics
  parameters      mean       std       mcse     ess_bulk
ess_tail      rha ...
  Symbol  Float64  Float64  Float64  Float64 ...
  Float64  Float64  ...
  695.4015      0.999 ...
  2 columns omitted

Quantiles
  parameters      2.5%     25.0%     50.0%     75.0%
97.5%
  Symbol  Float64  Float64  Float64  Float64 ...
  Float64
  p      0.3355    0.3954    0.4333    0.4673
  0.5412

```

4.1.1 The posterior of p is max at \$ p=0.4311\$, almost same as MLE. Not 0.5!

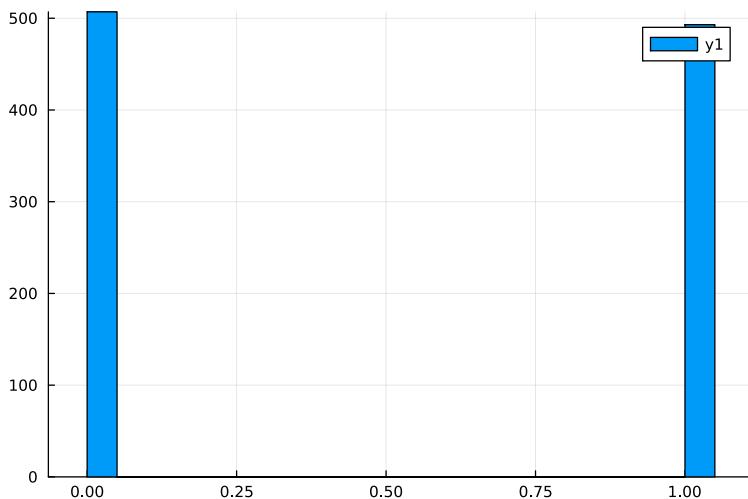
Because:

1. the data is 43% up;
2. we used a flat prior (uniform $\in [0, 1]$).

```
1 md"## 4.1.1 The posterior of p is max at $ p=0.4311$,
2 almost same as MLE. Not 0.5!
3 Because:
4 1) the data is 43% up;
5 2) we used a flat prior (uniform $\in [0,1]$). "
```

4.2 Input is 10X more data.

```
1 md"## 4.2 Input is 10X more data."
```

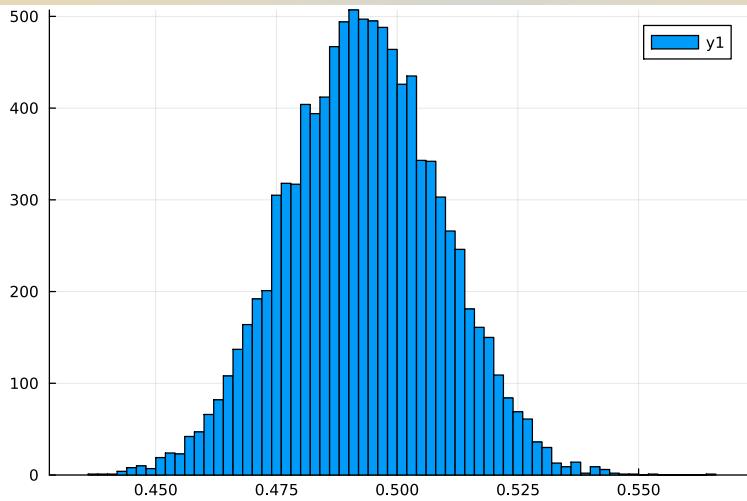


```
1 begin
2 # Fix the seed and draw data from a Bernoulli
3 distribution, i.e. draw heads or tails.
4 Random.seed!(7)
5 data2 = rand(Bernoulli(p_true), 1000)
6 println("The number of ups is
7 $(sum(data2))/$(length(data2)) ")
histogram(data2, bins=20)
end
```

```
The number of ups is 493/1000
```

4.2.1 sample using the same model but with 10X iterations

```
1 md"## 4.2.1 sample using the same model but with 10X
iterations"
```



```

1 begin
2   # sampling using the same model but more iterations
3   #<math>\epsilon=0.05</math>, <math>\tau=10</math>
4   @time chain_10X_HMC = sample(coinflip(data2), HMC(0.05,
5   10), 10_000)
6   display(chain_10X_HMC)
7   # Plot a summary of the sampling process for the
8   # parameter <math>p</math>, i.e. the probability of heads in a coin.
9   histogram(chain_10X_HMC[:p])
end

```

100%

```

3.998461 seconds (5.17 M allocations: 274.212 MiB, 2.04% gc time)
Chains MCMC chain (10000×11×1 Array{Float64, 3}):

Iterations       = 1:1:10000
Number of chains = 1
Samples per chain = 10000
Wall duration    = 3.93 seconds
Compute duration = 3.93 seconds
parameters       = p
internals        = lp, n_steps, is_accept, acceptance_rate, log_density, hamiltonian_energy, hamiltonian_energy_error, numerical_error, step_size, nom_step_size

Summary Statistics
  parameters      mean       std       mcse     ess_burn
k  ess_tail      ...
  Symbol  Float64  Float64  Float64  Float64
4  Float64  Fl ...
  p      0.4930   0.0159   0.0001  15496.752
1  10138.5250   1 ...
  2 columns omitted

Quantiles
  parameters    2.5%     25.0%    50.0%    75.0%
97.5%
  Symbol  Float64  Float64  Float64  Float64
Float64
  p      0.4619   0.4821   0.4930   0.5037
0.5240

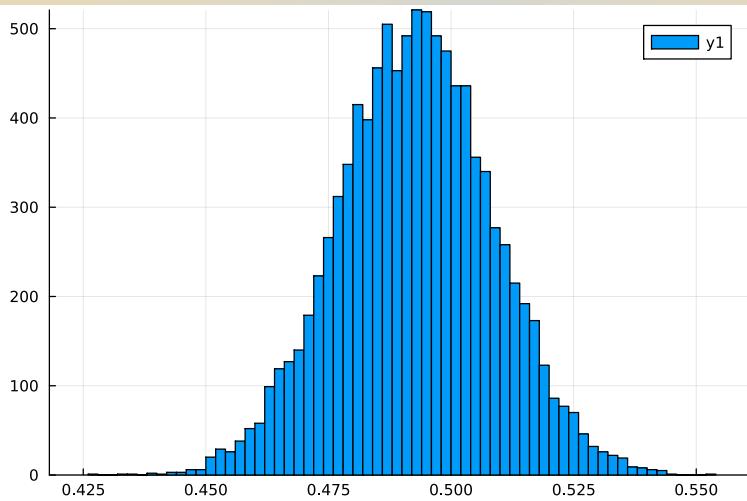
```

4.2.2 same model, but using the NUTS sampler (default parameters)

```

1 md"## 4.2.2 same model, but using the NUTS sampler
  (default parameters)"

```



```

1 begin
2   # Start sampling.
3   @time chain_10X_NUTS = sample(coinflip(data2), NUTS(),
4   10_000)
5   display(chain_10X_NUTS)
6   # Plot a summary of the sampling process for the
7   # parameter p, i.e. the probability of heads in a coin.
8   histogram(chain_10X_NUTS[:p])
end

```

100%

① Found initial step size
ε: 0.2

```

6.295829 seconds (4.19 M allocations: 276.446 MiB, 1.45% gc time)
Chains MCMC chain (10000×13×1 Array{Float64, 3}):

Iterations       = 1001:1:11000
Number of chains = 1
Samples per chain = 10000
Wall duration    = 6.18 seconds
Compute duration = 6.18 seconds
parameters       = p
internals        = lp, n_steps, is_accept, acceptance_rate,
log_density, hamiltonian_energy, hamiltonian_energy_error,
max_hamiltonian_energy_error, tree_depth,
numerical_error, step_size, nom_step_size

Summary Statistics
  parameters      mean       std       mcse     ess_bulk
ess_tail      r ...
  Symbol  Float64  Float64  Float64  Float64
Float64  Float ...
  p      0.4928    0.0157    0.0002  4407.9601
6543.2979  1.0 ...
  2 columns omitted

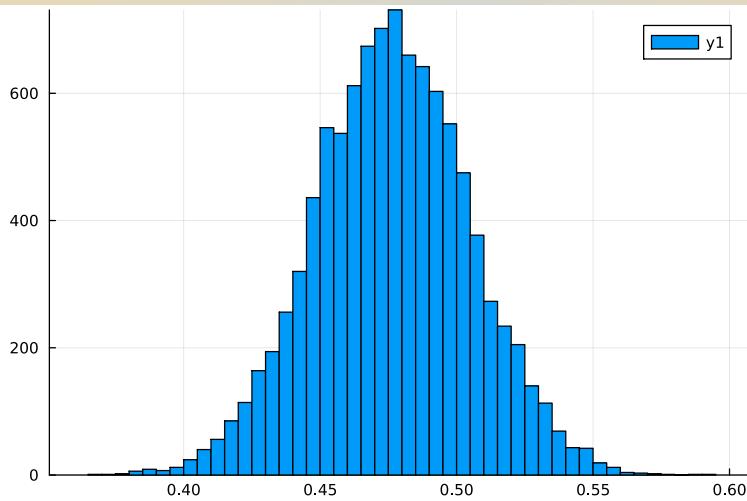
Quantiles
  parameters      2.5%     25.0%     50.0%     75.0%
97.5%
  Symbol  Float64  Float64  Float64  Float64
Float64
  p      0.4622    0.4821    0.4928    0.5033
0.5238

```

4.3 Modify the prior to alter the posterior

```
1 md"## 4.3 Modify the prior to alter the posterior"
```

```
coinflip_prior_around_0_5 (generic function with 2 methods)
1 # Declare our Bayesian model, p ~ beta, x ~ Bernoulli(p)
2 # But let's modify the \alpha and \beta of the prior
3 distribution.
4
5 @model function coinflip_prior_around_0_5(x)
6     # Our prior belief about the probability of heads in a
7 coin.
8     p ~ Beta(100, 100)
9
10    # The number of observations.
11    N = length(x)
12    for n in 1:N
13        # Heads or tails of a coin are drawn from a
14 Bernoulli distribution.
15        x[n] ~ Bernoulli(p)
16    end
17 end
```



```

1 begin
2   # Start sampling.
3   @time chain_strong_beta_NUTS =
4   sample(coinflip_prior_around_0_5(data), NUTS(), 10_000)
5   display(chain_strong_beta_NUTS)
6   # Plot a summary of the sampling process for the
7   # parameter p, i.e. the probability of heads in a coin.
7   histogram(chain_strong_beta_NUTS[:p])
end

```

100%

① Found initial step size
ε: 0.4

```

4.004209 seconds (5.56 M allocations: 373.837 MiB, 3.37% gc time, 53.42% compilation time)
Chains MCMC chain (10000×13×1 Array{Float64, 3}):

Iterations          = 1001:1:11000
Number of chains    = 1
Samples per chain   = 10000
Wall duration       = 3.03 seconds
Compute duration    = 3.03 seconds
parameters          = p
internals           = lp, n_steps, is_accept, acceptance_rate, log_density, hamiltonian_energy, hamiltonian_energy_error, max_hamiltonian_energy_error, tree_depth, numerical_error, step_size, nom_step_size

Summary Statistics
  parameters      mean       std       mcse     ess_bulk
ess_tail        r ...
  Symbol  Float64  Float64  Float64  Float64
  Float64  Float64 ...
  p      0.4765  0.0284  0.0004  4071.7624
6847.0813  1.0 ...
  2 columns omitted

Quantiles
  parameters    2.5%     25.0%     50.0%     75.0%
97.5%
  Symbol  Float64  Float64  Float64  Float64
  Float64 ...
  p      0.4203  0.4571  0.4764  0.4955
0.5331

```

4.3.1 Conclusion

- With a Strong prior beta(100,100) around 0.5, we can obtain a posterior estimate around 0.48, instead of 0.43.

```
1 md"## 4.3.1 Conclusion
2 - With a Strong prior beta(100,100) around 0.5, we can
   obtain a posterior estimate around 0.48, instead of 0.43."
```

5 Test two different Negative Binomial models

1. with a beta prior for p and Uniform prior for r
2. with a beta prior for p and Gamma prior for r

```
1 md"# 5 Test two different Negative Binomial models
2 1. with a beta prior for p and Uniform prior for r
3 1. with a beta prior for p and Gamma prior for r "
```

```
negative_binomial_model (generic function with 2 methods)
1 begin
2     # Generate synthetic data with known parameters
3     Random.seed!(28)
4     n = 1000
5     r = 3
6     p = 0.2
7     y = rand(NegativeBinomial(r, p), n)
8
9     # Bayesian Negative binomial model
10    @model nb_bayes(y) = begin
11        # Defining variables
12        r ~ Uniform(0, 100)
13        p ~ Beta(2, 2)
14
15        # Likelihood function
16        y_dist = NegativeBinomial(r, p)
17        y ~ y_dist
18    end
19
20
21    # Define the negative binomial likelihood
22    @model function negative_binomial_model(y)
23        # Prior distribution for the success probability
24        p ~ Beta(1, 1)
25        # Prior distribution for the dispersion parameter
26        r ~ Gamma(1, 1)
27        y_dist = NegativeBinomial(r, p)
28        # Likelihood
29        for i in eachindex(y)
30            y[i] ~ y_dist
31        end
32    end
33
34 end
```

```
► (1000)
1 size(y)
```

5.1 NB with a beta prior for p and Uniform prior for r

```
1 md"## 5.1 NB with a beta prior for p and Uniform prior for
r"
```

```

1 begin
2   # Sampling from the posterior distribution
3   @time chain_w_uniform_and_beta_prior =
4     sample(nb_bayes(y), NUTS(), 1_000)
5   display(chain_w_uniform_and_beta_prior)
end

```

100%

① Found initial step size
 ϵ : 0.0125

```

11.588600 seconds (8.27 M allocations: 571.869 M
iB, 2.20% gc time, 63.28% compilation time)
Chains MCMC chain (1000x14x1 Array{Float64, 3}):

Iterations       = 501:1:1500
Number of chains = 1
Samples per chain = 1000
Wall duration    = 9.76 seconds
Compute duration = 9.76 seconds
parameters       = r, p
internals        = lp, n_steps, is_accept, acceptance_rate,
log_density, hamiltonian_energy, hamiltonian_energy_error,
max_hamiltonian_energy_error, tree_depth, numerical_error, step_size, nominal_step_size

Summary Statistics
  parameters      mean       std       mcse      ess_bulk
k  ess_tail      rha ...
Symbol  Float64  Float64  Float64  Float64
4  Float64  Float64  ...
r      3.0357    0.1669    0.0109    260.173
6  248.9483   1.005 ...
p      0.1988    0.0092    0.0006    253.283
2  276.1358   1.004 ...
2 columns omitted

Quantiles
  parameters    2.5%     25.0%     50.0%     75.0%
97.5%
Symbol  Float64  Float64  Float64  Float64
Float64
r      2.7451    2.9313    3.0316    3.1235
3.3961
p      0.1824    0.1929    0.1990    0.2040

```

	parameters	mean	std	mcse	ess_bulk
1	:r	3.03573	0.166911	0.0108527	260.174
2	:p	0.198843	0.00922964	0.000609783	253.283

```

1 # Summarizing posterior distribution
2 @time describe(chain_w_uniform_and_beta_prior)

```

```

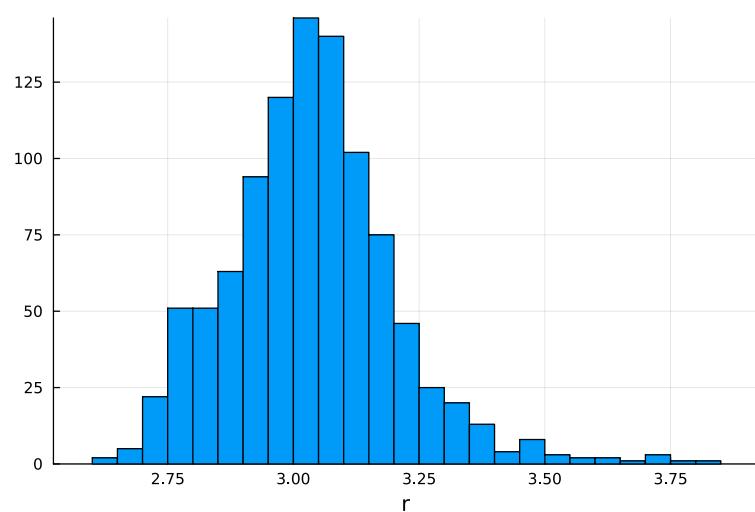
0.001812 seconds (916 allocations: 909.609 KiB)  ?

```

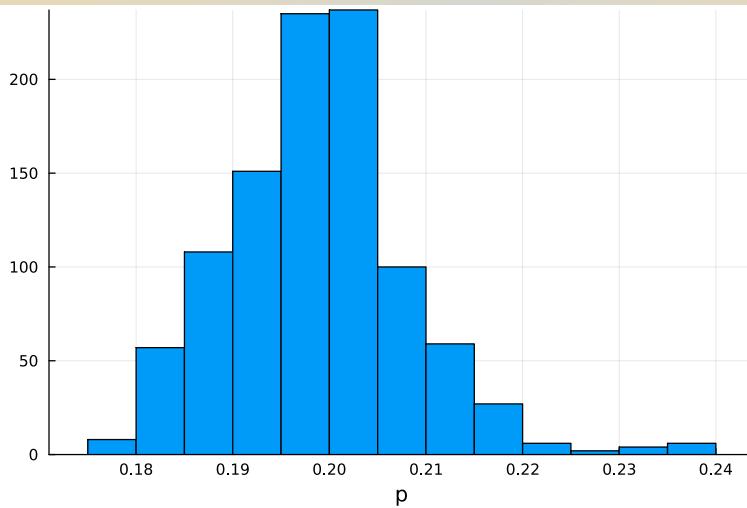
variable	mean	min	median	ma
1 :iteration	1000.5	501	1000.5	1500
2 :chain	1.0	1	1.0	1
3 :r	3.03573	2.61791	3.03156	3.809
4 :p	0.198843	0.176439	0.198989	0.238
5 :lp	-3363.98	-3371.6	-3363.65	-3363.65
6 :n_steps	7.178	1.0	7.0	31.0
7 :is_accept	1.0	1.0	1.0	1.0
8 :acceptance_rate	0.804302	0.000941705	0.902608	1.0
9 :log_density	-3363.98	-3371.6	-3363.65	-3363.65
10 :hamiltonian_energy	3364.94	3363.03	3364.67	3371.
: more				
16 :nom_step_size	0.346126	0.346126	0.346126	0.346126

```
1 @time describe(DataFrame(chain_w_uniform_and_beta_prior))
```

0.411745 seconds (215.26 k allocations: 15.133 MiB, 99.63% compilation time: 61% of which was recompilation)



```
1 histogram(chain_w_uniform_and_beta_prior[:r], xlabel="r",
legend=false)
```



```
1 histogram(chain_w_uniform_and_beta_prior[:p], xlabel="p",
legend=false)
```

5.2 NB with a beta prior for p and Gamma prior for r

```
1 md"## 5.2 NB with a beta prior for p and Gamma prior for r"
```

```
@time chain_with_gamma_and_beta_prior =
```

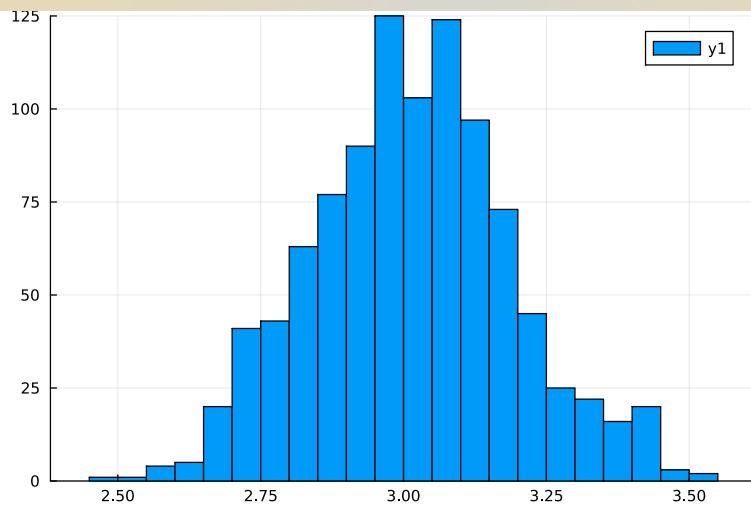
	iteration	chain	p	r	lp	n_steps	is_a
1	501	1	0.208408	3.13897	-3362.46	7.0	1.0
2	502	1	0.199412	3.09801	-3361.88	7.0	1.0
3	503	1	0.188981	2.85574	-3361.71	15.0	1.0
4	504	1	0.201038	3.06947	-3361.43	9.0	1.0
5	505	1	0.200822	3.05335	-3361.43	3.0	1.0
6	506	1	0.192431	2.92301	-3361.44	11.0	1.0
7	507	1	0.202756	3.11294	-3361.57	9.0	1.0
8	508	1	0.185653	2.79	-3362.15	11.0	1.0
9	509	1	0.184202	2.77843	-3362.3	3.0	1.0
10	510	1	0.182481	2.76941	-3362.58	3.0	1.0
	⋮ more						

```
1 @time chain_with_gamma_and_beta_prior =
sample(negative_binomial_model(y), NUTS(0.65), 1000)
```

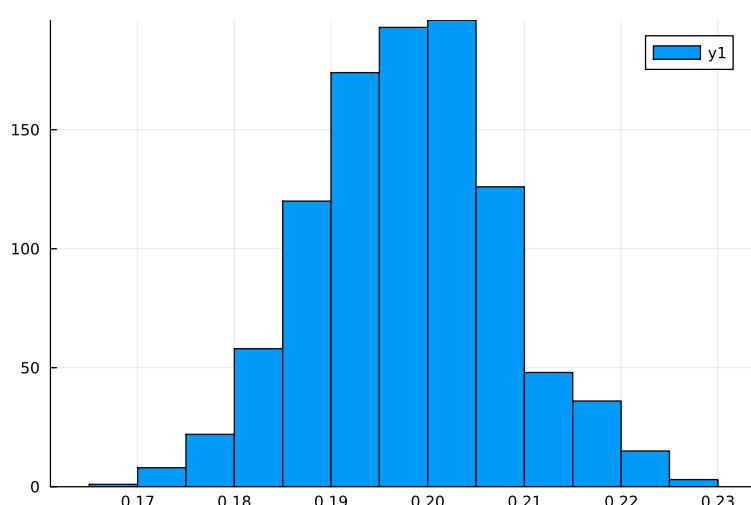
100%

① Found initial step size
ε: 0.025

10.284404 seconds (4.91 M allocations: 343.143 MiB, 1.86% gc time, 48.82% compilation time) ②



```
1 histogram(chain_with_gamma_and_beta_prior[:r])
```



```
1 histogram(chain_with_gamma_and_beta_prior[:p])
```

6 Estimate p of Binomial(N, p) from Binomial data with different Ns

```
1 md"# 6 Estimate p of Binomial(N, p) from Binomial data with  
different Ns"
```

```
simulate_data (generic function with 4 methods)  
1 function simulate_data(p::Float64, cov_mean::Int=4000,  
cov_std::Int=700, no_of_samples::Int=500)  
2     # Generate synthetic data with known parameters  
3     # cov_mean, cov_std are the parameters for the normal  
4 distribution.  
5  
6     # fix the seed  
7     Random.seed!(10)  
8     # Create a normal distribution centered around 4000  
9     cov_dist = Normal(cov_mean, cov_std)  
10  
11    # Sample an integer coverage vector from the  
12 distribution  
13    N_v = round.(rand(cov_dist, no_of_samples))  
14  
15    x_v = zeros(Int, length(N_v))  
16    for (i, N) in enumerate(N_v)  
17        x_v[i] = rand(Binomial(N, p))  
18    end  
19    vaf_v = x_v ./ N_v  
20    return N_v, x_v, vaf_v  
end
```

```
1 begin  
2     @time N_v, x_v, vaf_v = simulate_data(0.0005)  
3     @time N_v_005, x_v_005, vaf_v_005 = simulate_data(0.005)  
4     @time N_v_00001, x_v_00001, vaf_v_00001 =  
5         simulate_data(0.00001)  
6  
7     @time println("The number of samples with VAF=0:  
8         $(sum(vaf_v_00001==0) .);  
end
```

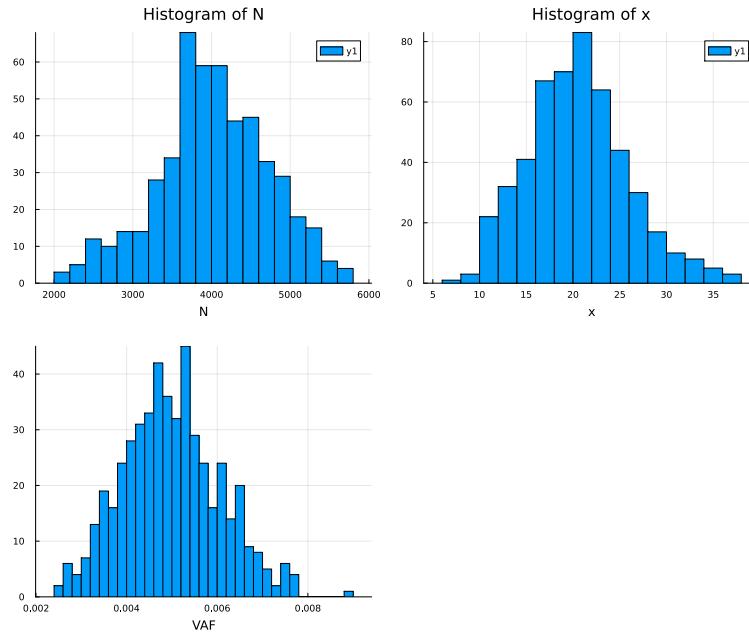
```
0.000094 seconds (10 allocations: 16.703 KiB)  
0.000097 seconds (10 allocations: 16.703 KiB)  
0.000038 seconds (10 allocations: 16.703 KiB)  
The number of samples with VAF=0: 490.  
0.139649 seconds (122.03 k allocations: 8.158 MiB, 1  
8.22% gc time, 99.79% compilation time)
```

```
Vector{Float64} (alias for Array{Float64, 1})  
1 typeof(vaf_v)
```

```

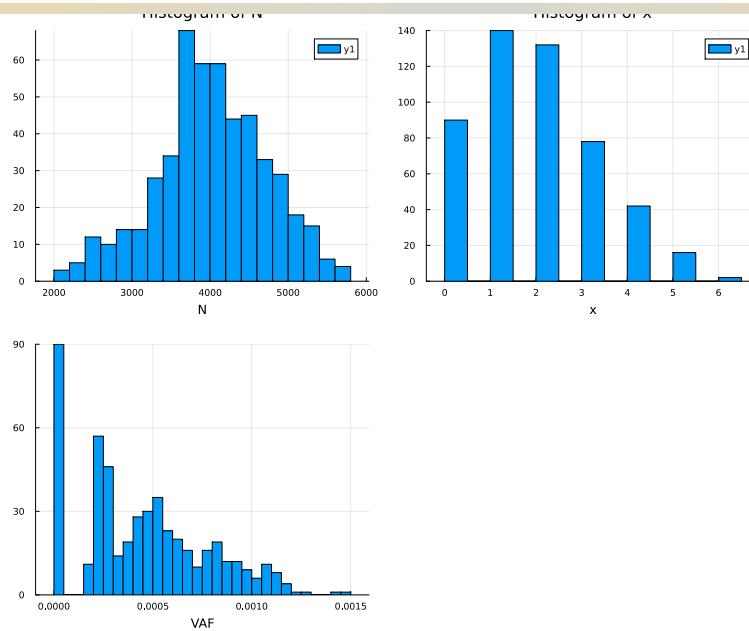
plot_histogram (generic function with 1 method)
1 function plot_histogram(N_v::Vector{Float64},
2   x_v::Vector{Int64}, vaf_v::Vector{Float64}, p::Float64)
3   @time h1 = histogram(N_v, title="Histogram of N",
4   xaxis="N")
5   @time h2 = histogram(x_v, title="Histogram of x",
6   xaxis="x")
7   @time h_of_vaf_v = histogram(vaf_v, xaxis="VAF",
7 bins=30, legend=false)
8   @time plot(h1, h2, h_of_vaf_v, layout=(2,2),
9   left_margin=5*Plots.mm,
10  bottom_margin=5*Plots.mm, size=(1000,830) )
end

```



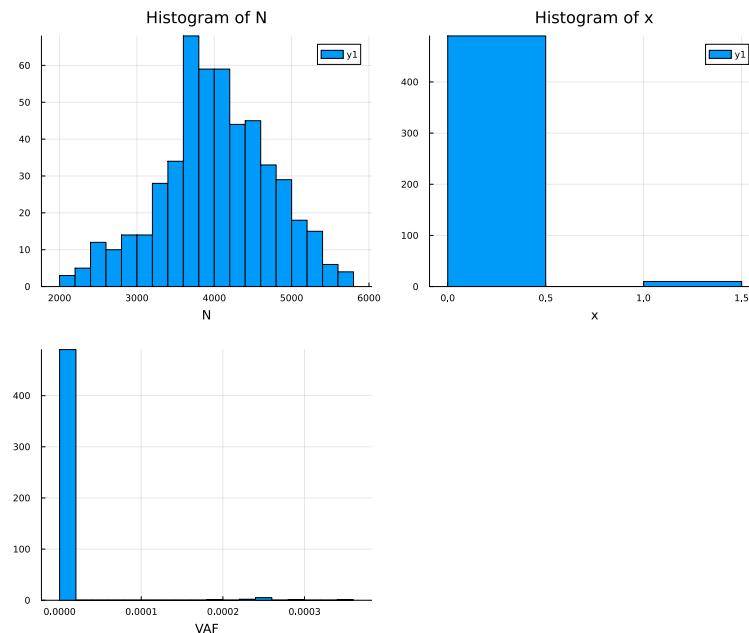
```
1 plot_histogram(N_v_005, x_v_005, vaf_v_005, 0.005)
```

B) 0.001736 seconds (4.23 k allocations: 234.562 KiB) ②
 0.001332 seconds (4.20 k allocations: 234.062 KiB)
 0.001255 seconds (4.42 k allocations: 241.750 KiB)
 0.135867 seconds (169.07 k allocations: 11.719 MiB,
 98.34% compilation time)



```
1 plot_histogram(N_v, x_v, vaf_v, 0.0005)
```

```
0.001699 seconds (4.23 k allocations: 234.562 KiB)
B)
0.001146 seconds (4.15 k allocations: 232.406 KiB)
0.001224 seconds (4.38 k allocations: 240.594 KiB)
0.001369 seconds (1.49 k allocations: 301.781 KiB)
```



```
1 plot_histogram(N_v_00001, x_v_00001, vaf_v_00001, 0.00001)
```

```
0.001799 seconds (4.23 k allocations: 234.562 KiB)
B)
0.001162 seconds (4.02 k allocations: 225.906 KiB)
0.001190 seconds (4.22 k allocations: 225.031 KiB)
0.001459 seconds (1.49 k allocations: 303.516 KiB)
```

```

binomial_w_beta_prior (generic function with 2 methods)
1 # Define the model
2 @model binomial_w_beta_prior(N, x) = begin
3     # Prior distribution for the success probability
4     p ~ Beta(1, 5)
5     # Prior distribution for the dispersion parameter, for
6     # the NB distribution
7     # r ~ Gamma(1, 1)
8     #y_dist = NegativeBinomial(r, p)
9     # Likelihood
10    for i in eachindex(N)
11        x[i] ~ Binomial(N[i], p)
12    end
13 end

```

6.1 Estimate if true p is 0.0005

```
1 md"## 6.1 Estimate if true p is 0.0005"
```

```

1 begin
2     # Sampling from the posterior distribution
3     @time chain_binomial_w_beta_prior =
4         sample(binomial_w_beta_prior(N_v, x_v), NUTS(), 2000)
5     display(chain_binomial_w_beta_prior)
6 end

```

100%

① Found initial step size
 $\epsilon: 0.003125$

```

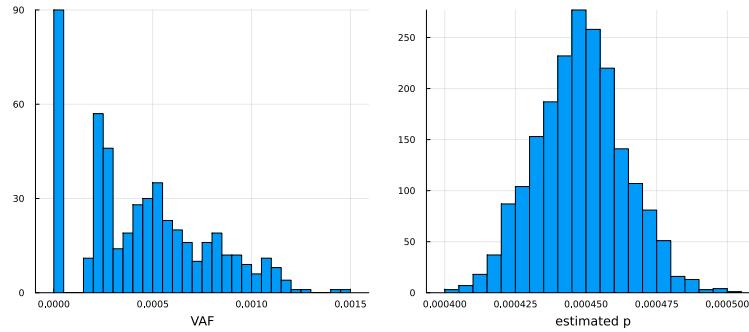
4.232009 seconds (2.62 M allocations: 180.172 MiB, 2.46% gc time, 51.77% compilation time)
Chains MCMC chain (2000×13×1 Array{Float64, 3}):
```

Iterations	1001:1:3000			
Number of chains	1			
Samples per chain	2000			
Wall duration	3.28 seconds			
Compute duration	3.28 seconds			
parameters	p			
internals	lp, n_steps, is_accept, acceptance_rate, log_density, hamiltonian_energy, hamiltonian_energy_error, max_hamiltonian_energy_error, tree_depth, numerical_error, step_size, nom_step_size			
Summary Statistics				
parameters	mean std mcse ess_bulk			
ess_tail	rha ...			
Symbol	Float64	Float64	Float64	Float64
Float64	Float64	...		
	0.0004	0.0000	0.0000	882.0261
999.6807	0.999 ...			
2 columns omitted				
Quantiles				
parameters	2.5%	25.0%	50.0%	75.0%
97.5%				
Symbol	Float64	Float64	Float64	Float64
Float64	...			
	0.0004	0.0004	0.0004	0.0005
0.0005				

```

plot_estimates (generic function with 1 method)
1 function plot_estimates(vaf_v, estimated_p)
2   h_of_vaf_v = histogram(vaf_v, xaxis="VAF", bins=30,
3   legend=false)
4   h_of_p_bin = histogram(estimated_p, xaxis="estimated
p", legend=false)
5   plot(h_of_vaf_v, h_of_p_bin, layout=(1,2),
left_margin=5*Plots.mm, bottom_margin=5*Plots.mm, size=
(1000,430) )
end

```



```
1 plot_estimates(vaf_v, chain_binomial_w_beta_prior[:p])
```

6.2 Estimate if true p is 0.005

```
1 md"## 6.2 Estimate if true p is 0.005"
```

```

1 begin
2     # Sampling from the posterior distribution
3     @time chain_binomial_w_beta_prior_005 =
4         sample(binomial_w_beta_prior(N_v_005, x_v_005), NUTS(),
5             2000)
6     display(chain_binomial_w_beta_prior_005)
end

```

100%

① Found initial step size
 ϵ : 0.003125

```

7.797897 seconds (1.14 M allocations: 76.075 MiB, 0.27% gc time)
Chains MCMC chain (2000x13x1 Array{Float64, 3}):
```

Iterations = 1001:1:3000
Number of chains = 1
Samples per chain = 2000
Wall duration = 7.78 seconds
Compute duration = 7.78 seconds
parameters = p
internals = lp, n_steps, is_accept, acceptance_rate, log_density, hamiltonian_energy, hamiltonian_energy_error, max_hamiltonian_energy_error, tree_depth, numerical_error, step_size, nom_step_size

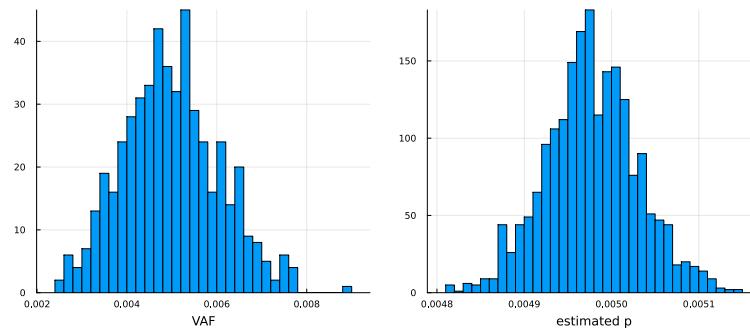
Summary Statistics

parameters	mean	std	mcse	ess_bulk
ess_tail	rh ...			
Symbol	Float64	Float64	Float64	Float64
Float64	Float64	Float64	Float64	Float64
p	0.0050	0.0001	0.0000	1589.6704
938.5041	0.99 ...			

2 columns omitted

Quantiles

parameters	2.5%	25.0%	50.0%	75.0%
97.5%				
Symbol	Float64	Float64	Float64	Float64
Float64	Float64	Float64	Float64	Float64
p	0.0049	0.0049	0.0050	0.0050
0.0051				



```

1 plot_estimates(vaf_v_005,
chain_binomial_w_beta_prior_005[:p])

```

6.3 Estimate if true p is 0.00001

```
1 begin
2     # Sampling from the posterior distribution
3     @time chain_binomial_w_beta_prior_00001 =
4         sample(binomial_w_beta_prior(N_v_00001, x_v_00001),
5             NUTS(), 2000)
6     display(chain_binomial_w_beta_prior_00001)
end
```

100%

① Found initial step size
ε: 0.003125

```
2.666185 seconds (1.22 M allocations: 80.550 MiB, 0.64% gc time)
Chains MCMC chain (2000x13x1 Array{Float64, 3}):
```

Iterations = 1001:1:3000
Number of chains = 1
Samples per chain = 2000
Wall duration = 2.64 seconds
Compute duration = 2.64 seconds
parameters = p
internals = lp, n_steps, is_accept, acceptance_rate, log_density, hamiltonian_energy, hamiltonian_energy_error, max_hamiltonian_energy_error, tree_depth, numerical_error, step_size, nom_step_size

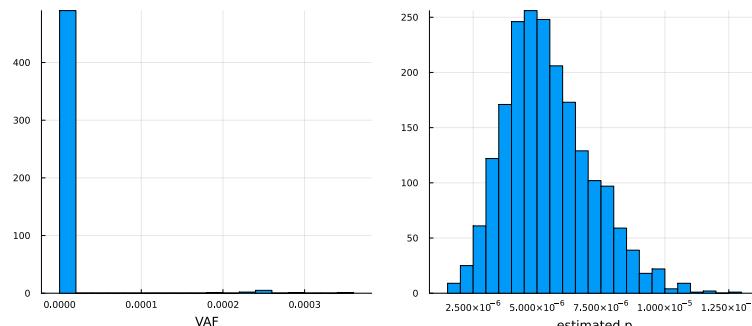
Summary Statistics

parameters	mean	std	mcse	ess_bulk
ess_tail	rh ...	Symbol	Float64	Float64
Float64	Float64	Float64	Float64	Float64
p	0.0000	0.0000	0.0000	821.2807
1048.3478	1.00 ...			

2 columns omitted

Quantiles

parameters	2.5%	25.0%	50.0%	75.0%
97.5%	Symbol	Float64	Float64	Float64
Float64	Float64	Float64	Float64	Float64
p	0.0000	0.0000	0.0000	0.0000
0.0000				



```
1 plot_estimates(vaf_v_00001,
chain_binomial_w_beta_prior_00001[:p])
```

7 Estimate p of NegativeBinomial(r, p) from Binomial data with different Ns

```
1 md"# 7 Estimate p of NegativeBinomial(r, p) from Binomial  
data with different Ns"
```

```
negbin_w_beta_prior (generic function with 2 methods)  
1 # Define the model  
2 @model negbin_w_beta_prior(N, x) = begin  
3     # Prior distribution for the success probability  
4     p ~ Beta(1, 5)  
5     # Prior distribution for the dispersion parameter, for  
6     # the NB distribution  
7     # r ~ Gamma(1, 1)  
8     #y_dist = NegativeBinomial(r, p)  
9     # Likelihood  
10    for i in eachindex(N)  
11        x[i] ~ NegativeBinomial(N[i]-x[i], p)  
12    end  
end
```

```

1 begin
2   # Sampling from the posterior distribution
3   @time chain_negbin_w_beta_prior =
4   sample(negbin_w_beta_prior(N_v, x_v), NUTS(), 2000)
5   display(chain_negbin_w_beta_prior)
end

```

100%

① Found initial step size
 $\epsilon: 0.003125$

```

5.385602 seconds (2.58 M allocations: 177.628 MiB, 1.31% gc time, 38.10% compilation time)
Chains MCMC chain (2000x13x1 Array{Float64, 3}):

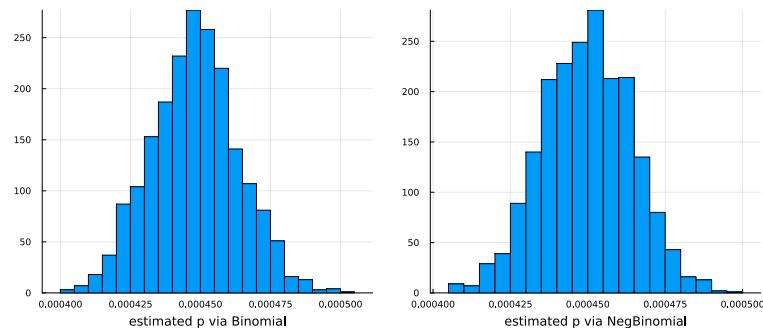
Iterations       = 1001:1:3000
Number of chains = 1
Samples per chain = 2000
Wall duration    = 4.52 seconds
Compute duration = 4.52 seconds
parameters       = p
internals        = lp, n_steps, is_accept, acceptance_rate, log_density, hamiltonian_energy, hamiltonian_energy_error, max_hamiltonian_energy_error, tree_depth, numerical_error, step_size, nom_step_size

Summary Statistics
  parameters      mean      std      mcse      ess_bulk
ess_tail      rh ...
  Symbol  Float64  Float64  Float64  Float64
Float64  Float64 ...
  p      0.9996  0.0000  0.0000  888.3702
1276.0136  1.00 ...

2 columns omitted

Quantiles
  parameters      2.5%      25.0%      50.0%      75.0%
97.5%
  Symbol  Float64  Float64  Float64  Float64
Float64 ...
  p      0.9995  0.9995  0.9996  0.9996
0.9996

```



```

1 begin
2   h_of_p_bin = histogram(chain_binomial_w_beta_prior[:p],
3     xaxis="estimated p via Binomial", legend=false)
4   h_of_p_negbin = histogram(
5     (chain_negbin_w_beta_prior[:p].-1),
6     xaxis="estimated p via NegBinomial", legend=false)
7   plot(h_of_p_bin, h_of_p_negbin, layout=(1,2),
8     left_margin=5*Plots.mm,
9     bottom_margin=5*Plots.mm, size=(1000,430) )
end

```

7.1 \hat{p} of Beta+Binomial is similar to \hat{p} of Beta+NegativeBinomial.

```
1 md"## 7.1  $\hat{p}$  of Beta+Binomial is similar to  $\hat{p}$  of Beta+NegativeBinomial."
```

```
"  $\hat{p}$  of Beta+Binomial, 0.000448 is similar to  $\hat{p}$  of Beta+NegativeBi
1 begin
2    $\hat{p}_1$  = median(chain_binomial_w_beta_prior[:p])
3    $\hat{p}_2$  = median(-(chain_negbin_w_beta_prior[:p].-1))
4   "  $\hat{p}$  of Beta+Binomial, $(@sprintf("%0.6f",  $\hat{p}_1$ )) is
5     similar to  $\hat{p}$  of Beta+NegativeBinomial,
6     $(@sprintf("%0.6f",  $\hat{p}_2$ ))."
7 end
```

8 Multinomial

```
1 md"# 8 Multinomial"
```

```
4x1000 Matrix{Int64}:
 4      3      5      4      0      4 ...      3      4      5
 15     14     11     10     19     17 ...     17     17     17
 6      3      4      3      6      5      5      5      3      5
39975  39980  39980  39983  39975  39974 ... 39975  39976  39973
```

```
1 begin
2   Random.seed!(250)
3   # Generate some dummy data
4   data_mn = rand(Multinomial(40000, [0.0001, 0.0004,
5   0.0001, 0.9994]), 1000)
6 end
```

```
multinomial_model (generic function with 2 methods)
```

```
1 # Define the model
2 @model function multinomial_model(data)
3   no_of_alleles, no_of_samples = size(data)
4
5   # Prior for the parameters (Dirichlet distribution)
6   α = ones(no_of_alleles) # Prior concentration
7   parameters
8
9   # Parameters (probabilities of the categories)
10  ε ~ Dirichlet(α)
11  # Likelihood (multinomial distribution)
12  for j in 1:no_of_samples
13    data[:, j] ~ Multinomial(sum(data[:,j]), ε)
14  end
15 end
```

```
"MCMCChains.Chains{Float64, AxisArrays.AxisArray{Float64, 3, Arra
```

```
1 begin
2   # Run the model
3   @time chain_mn = sample(multinomial_model(data_mn),
4   NUTS(), 1_000)
5   # Summarize the results
6   summary(chain_mn)
7 end
```

```
100%
```

ⓘ Found initial step size
ε: 0.000390625

9.936392 seconds (13.98 M allocations: 1.099 GiB, 3.18% gc time, 59.30% compilation time)

	iteration	chain	$\epsilon[1]$	$\epsilon[2]$	$\epsilon[3]$	⋮
1	501	1	0.000100408	0.000401299	9.60591e-5	0.9
2	502	1	9.79426e-5	0.000400827	9.88822e-5	0.9
3	503	1	0.00010151	0.000408165	0.000100396	0.9
4	504	1	0.000102396	0.000406606	0.000100826	0.9
5	505	1	0.000101313	0.000402725	9.85483e-5	0.9
6	506	1	0.000102202	0.000407345	0.000101775	0.9
7	507	1	9.99532e-5	0.000405024	9.72597e-5	0.9
8	508	1	0.000103476	0.000400731	9.62322e-5	0.9
9	509	1	0.000100917	0.000397146	9.82294e-5	0.9
10	510	1	9.71677e-5	0.000411256	0.000100168	0.9
⋮ more						

```
◀ ▶
1 chain_mn
```

```
▶ (:value, :logevidence, :name_map, :info)
1 propertynames(chain_mn)
```

```
3-dimensional AxisArray{Float64,3,...} with axes:
  :iter, 501:1:1500
  :var, [Symbol("ε[1]"), Symbol("ε[2]"), Symbol("ε[3]"), Symbol
  :chain, 1:1
And data, a 1000×16×1 Array{Float64, 3}:
[:, :, 1] =
0.000100408 0.000401299 9.60591e-5 ... 1.2582 2.0 0.0
9.79426e-5 0.000400827 9.88822e-5 ... -0.808108 2.0 0.0
0.00010151 0.000408165 0.000100396 ... -0.746499 2.0 0.0
0.000102396 0.000406606 0.000100826 ... -0.193956 2.0 0.0
0.000101313 0.000402725 9.85483e-5 ... 0.0763846 2.0 0.0
0.000102202 0.000407345 0.000101775 ... -0.160806 2.0 0.0
9.99532e-5 0.000405024 9.72597e-5 ... 0.657073 2.0 0.0
⋮
0.000102728 0.000403545 0.000101307 ... 0.309617 2.0 0.0
0.000100568 0.0004038 9.85621e-5 ... 0.062498 2.0 0.0
9.87572e-5 0.000406329 0.000102453 ... 0.61542 2.0 0.0
0.000104567 0.00040759 9.93938e-5 ... 0.66652 2.0 0.0
9.89762e-5 0.000402849 0.000102203 ... 0.21905 2.0 0.0
0.000105195 0.000403377 0.000100183 ... 0.67073 2.0 0.0
```

```
◀ ▶
1 chain_mn.value
```

missing

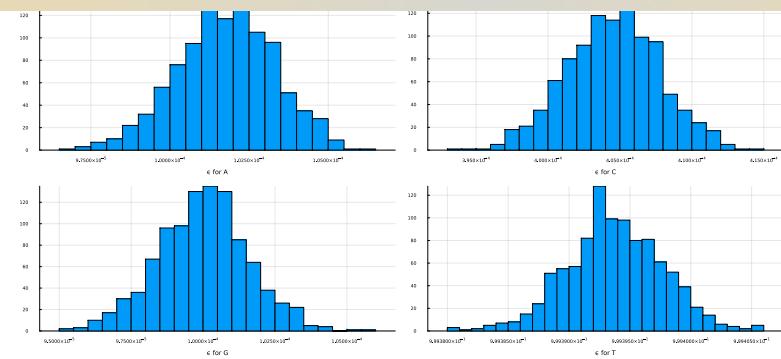
```
1 chain_mn.logevidence
```

```
▶ (varname_to_symbol = OrderedDict(ε[1] ⇒ Symbol("ε[1]"), ε[2] ⇒
```

```
◀ ▶
1 chain_mn.info
```

```
▶ (parameters = [Symbol("ε[1]"), Symbol("ε[2]"), Symbol("ε[3]"), ε
```

```
◀ ▶
1 chain_mn.name_map
```



```

1 begin
2   h_of_p1 = histogram(chain_mn[:"ε[1]"], xaxis="ε for A",
3   legend=false)
4   h_of_p2 = histogram(chain_mn[:"ε[2]"], xaxis="ε for C",
5   legend=false)
6   h_of_p3 = histogram(chain_mn[:"ε[3]"], xaxis="ε for G",
7   legend=false)
8   h_of_p4 = histogram(chain_mn[:"ε[4]"], xaxis="ε for T",
9   legend=false)

  plot(h_of_p1, h_of_p2, h_of_p3, h_of_p4, layout=(2,2),
       left_margin=5*Plots.mm,
       bottom_margin=10*Plots.mm, size=(2000,930) )
end

```

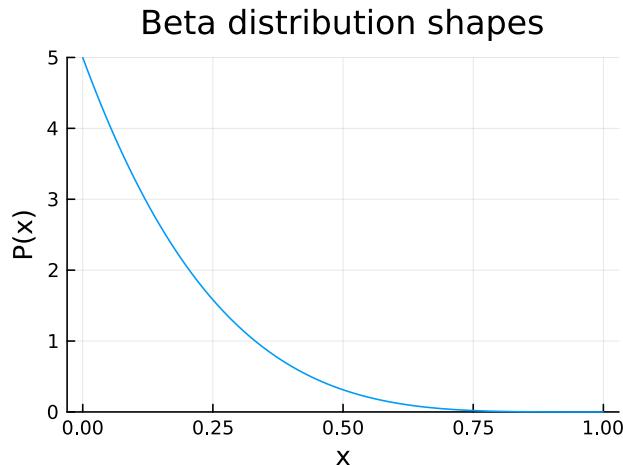
```

1 using Plots
2
3 # Plot the cluster assignments over time

```

10 Visualize beta distributions

```
1 md "# 10 Visualize beta distributions"
```



```
1 plot(x -> pdf(Beta(1, 5), x),
2     minimum(support(Beta(1,1))):0.01:maximum(support(Beta(1,
3     1))),
4     ylim=(0, 5), size=(400, 300), label=false,
5     xlabel="x", ylabel="P(x)", title="Beta distribution
shapes")
```

```
0.0:0.01:1.0
```

```
1 0:0.01:1
```

```
#9 (generic function with 1 method)
```

```
1 x -> pdf(Beta(1, 5), x)
```